

ANKUSH PRAMANIK

✉ ankushpramanik@gmail.com | +91 9818807215 |  LinkedIn |  GitHub |  Portfolio | Ghaziabad, UP, India

SUMMARY

AI Engineer building applied GenAI systems — RAG pipelines, agentic AI frameworks, and LLM-integrated backend services. Core strengths: multi-provider LLM orchestration (OpenAI, Anthropic, Vertex AI), AI architecture design, and production-grade retrieval. Builds with Claude Code, Cursor and Codex, engineering reliable AI systems.

EXPERIENCE

Ginclair | *AI/ML Engineer · Full-time* Jul 2026 – Present | Bengaluru, India (Remote)

- Architected **agentic AI** features with GraphRAG, PydanticAI, CrewAI and OpenAI SDK, building **multi-agent orchestration** pipelines via **LangChain + LangGraph** in Python.
- Developed **RAG retrieval** systems using **FAISS vector search** and Gemini embeddings with FastAPI, integrating prompt engineering and **structured JSON** outputs for production LLM apps.
- Benchmarked AI architecture patterns (**ReAct loops, tool-use agents, guardrails**) across **OpenAI GPT**, Claude, and Gemini APIs, authoring internal technical documentation.

Search Precise Consultancy | *AI Engineer Intern* Jun 2026 – Aug 2026 | Gurugram, India

- Developed **prompt engineering** solutions using **Google Gemini** for technical diagram generation, system architecture, workflow visualization, and AI-assisted documentation.
- Designed structured prompts for **AI-driven patient-data diagnostic workflows** involving **deep learning, data processing, diagnostic inference, and interpretable outputs**.
- Optimized **LLM prompting and structured outputs** for improved accuracy, consistency, visual quality, and reliability of AI-generated technical artifacts.

Capgemini | *AI Java Backend Intern* Jan 2026 – May 2026 | Onsite

- Architected **RESTful microservices** with Java 17, Spring Boot, Spring Security (RBAC), JPA/Hibernate, and pgvector-backed PostgreSQL for production systems.
- Engineered RAG-based **LLM orchestration** pipelines with **Spring AI**, automating document data extraction via **structured output** validation.
- Containerized services with **Docker**/ Docker Compose, architected CI/CD pipelines via **GitHub Actions**, and validated API reliability through Postman integration testing.

PROJECTS

Jarvis – Agentic AI Workspace with Hybrid RAG & Memory Orchestration Demo Link |  Repo

Python | SQL | FastAPI | RESTful APIs | Pydantic | LLM | RAG | Embeddings | AI Agents | JSON | Docker

- Architected an **agentic AI workspace** with an **AI intent router** for **hybrid RAG**, web search, data analysis, memory, and utility tools using **structured JSON outputs** and self-verification.
- Built a **FastAPI** backend with **REST APIs**, SSE streaming, **JWT authentication**, per-user isolation, LLM/search integrations, **Pydantic validation**, guardrails, and a **JavaScript** voice-enabled frontend, containerized with **Docker**.

Code Explainer Enterprise – GenAI Code Analyzer Demo Link |  Repo

FastAPI | Pydantic | OpenAI SDK | uv | LangGraph | Docker | Langfuse

- Built a **GenAI code analysis** platform with **FastAPI** and **Pydantic**, using the **OpenAI SDK** and **LangGraph** to orchestrate structured code-analysis workflows across 6 programming languages.
- Containerized services with **Docker**, managed dependencies using **uv**, and integrated **Langfuse** for LLM tracing and observability across Big-O analysis, security detection, and automated refactoring pipelines.

TECHNICAL SKILLS

Languages:	Python, Java, SQL, JavaScript (Basic)
Generative AI / LLMs:	LLM, MCP, LangChain, Google ADK, Vertex AI, RAG, OpenAI, Google Gemini, Anthropic Claude, vLLM, Ollama, Prompt Engineering, Embeddings
Agentic AI / Architecture:	AI Agents, AI Architecture, Multi-Agent Orchestration, ReAct, Autonomous Agent Loops, PydanticAI, CrewAI, LangGraph
Retrieval / Vectors:	FAISS, Pinecone, Qdrant, ChromaDB, LlamaIndex, Haystack, Hybrid Search, GraphRAG
Data / Backend:	FastAPI, RESTful APIs, API Integration, PostgreSQL, MySQL, Supabase, Pydantic, JSON, Webhooks, JWT
Cloud / DevOps:	Google Cloud Platform (GCP), Vertex AI Platform, Docker, Microsoft Azure, Amazon Bedrock
Eval / Quality:	Output Validation, Guardrails, Structured Outputs, JSON Schema, RAGAS, LangSmith, Langfuse
Professional Skills:	Requirements Gathering, Solution Design, Cross-Functional Collaboration, Stakeholder Communication, Technical Documentation, System Design

EDUCATION

ABES Engineering College, Ghaziabad, UP Aug 2022 – May 2026

B.Tech – Computer Science & Engineering (Data Science Specialization) | **78.86%**

St. Teresa School, Ghaziabad, UP 2020 – 2022

Class 12 – CBSE: 78.3% | Class 10 – CBSE: 82%

CERTIFICATIONS

Introduction to Model Context Protocol – Anthropic (Jul 2026)	 Verify
Introduction to Agent Skills – Anthropic (Jul 2026)	 Verify
AI Fluency Framework & Foundations – Anthropic (Jul 2026)	 Verify